

## HW2: Command-line design

---

In this assignment, you'll incrementally build up a simple Python command-line tool. Please note that for this homework assignment, you will be graded on the final tool and your associated write-up; you only need to turn in intermediate steps if you are unable to complete step #3.

### Goal

While there are many formats used for tagging data, one of the most common is multi-column tabular format with a blank line between each sentence. For instance, here are two sentences as they appear in the CoNLL-2000 shared task on chunking (Tjong Kim Sang & Buchholz 2000):

```
The DT B-NP
August NNP I-NP
deficit NN I-NP
and CC 0
the DT B-NP
# # I-NP
2.2 CD I-NP
billion CD I-NP
gap NN I-NP
registered VBN B-VP
in IN B-PP
July NNP B-NP
are VBP B-VP
topped VBN I-VP
only RB B-ADVP
by IN B-PP
the DT B-NP
# # I-NP
2.3 CD I-NP
billion CD I-NP
deficit NN I-NP
of IN B-PP
October NNP B-NP
1988 CD I-NP
. . 0

Sanjay NNP B-NP
Joshi NNP I-NP
, , 0
European JJ B-NP
economist NN I-NP
at IN B-PP
```

```
Baring NNP B-NP
Brothers NNPS I-NP
& CC I-NP
Co. NNP I-NP
, , 0
said VBD B-VP
there EX B-NP
is VBZ B-VP
no DT B-NP
sign NN I-NP
that IN B-SBAR
Britain NNP B-NP
's POS B-NP
manufacturing NN I-NP
industry NN I-NP
is VBZ B-VP
transforming VBG I-VP
itself PRP B-NP
to TO B-VP
boost VB I-VP
exports NNS B-NP
. . 0
```

Here the first column is the token, the second is a hypothesized POS tag, and the third is the chunk tag. In Python we might represent

- each row as a list or tuple (here, we choose the former),
- each sentence as a list of such lists, and
- the entire corpus as a list of sentences.

In this assignment, you will incrementally develop a Python command-line tool which automatically reads in tagging data in the above format and randomly splits it into training, development, and test data, writing the resulting "splits" to new text files.

## Step 1: parsing

For data in the above format, we can read the data one sentence at a time using the following generator function.

```
from typing import Iterator, List

def read_tags(path: str) -> Iterator[List[List[str]]]:
    with open(path, "r") as source:
        lines = []
        for line in source:
```

```

        line = line.rstrip()
        if line: # Line is contentful.
            lines.append(line.split())
        else: # Line is blank.
            yield lines.copy()
            lines.clear()
# Just in case someone forgets to put a blank line at the end...
if lines:
    yield lines

```

Then, to read the entire corpus at once, then, we can do the following, for instance:

```
corpus = list(read_tags("conll2000.txt"))
```

## What to do

Using the `argparse` module, build a Python command-line tool which takes four arguments: `input`, `train`, `dev`, and `test`. The program should:

1. Read all of the `input` data in using the above snippet.
2. Split the data into an 80% training set, 10% development set, and 10% test set.
3. Write the training set to the `train` path.
4. Write the development set to the `dev` path.
5. Write the testing set to the `test` path.

The resulting training, development, and testing set files should be in the same columnar format as the input format.

## Hints

- Read the [argparse docs](#) before you get started.
- The skeleton of a program is provided in `split.py`.
- Do not hard-code any of the paths.
- You may have to mark your code executable, which you can do by running `chmod +x split.py` from the command-line.
- To stay [DRY](#), you may want to write a function `write_tags` which writes the data out.

## Testing

1. Confirm that the following works as expected:

```
$ ./split.py conll2000.tag train.tag dev.tag test.tag
```

2. Manually check that the file lengths are correct using the `wc` command-line tool:

```
$ wc -l train.tag dev.tag test.tag
```

Note that your exact lengths may differ slightly (because the sentence lengths differ slightly) but that `dev.tag` and `test.tag` should be roughly the same lengths, and both should be roughly 8x shorter than `train.tag`.

## Step 2: Randomization

Above, we did not specify how the data was split. Now, we will modify it so that the division is pseudo-random.

### What to do

Modify `split.py` from the previous step so that the split into training, development, and testing data is pseudo-random. The easiest way to do this is to reuse the code from the previous step, but randomly shuffle the list of sentences using `random.shuffle`.

### Hints

- Read the [random docs](#) before you get started.
- Note that `random.shuffle` **works in-place**.

## Step 3: Seeding

When Python loads the `random` module, it automatically "seeds" the pseudo-random number generator (PRNG) with the current wall clock time. Thus every time one runs `./split.py`, a different split is obtained. This makes it extremely difficult to exactly replicate a workflow that involves random number generation. One solution involves having the user provide a seed (an integer, usually) to the PRNG before it is used. So long as the user keeps a record of the seed used (or uses a memorable number, like a birthdate or an address) anyone can replicate the "random" behavior in the future.

### What to do

Add a `--seed` flag to your script, and use this seed to seed the PRNG before randomizing the data. Mark the flag mandatory in the `argparse` declaration so that a user must specify some seed.

### Testing

1. Confirm that the following succeeds:

```
$ ./split.py --seed=272 conll2000.tag train.tag dev.tag test.tag
```

2. Confirm that your script prints out an informative error message if the user omits the `--seed` flag. E.g.:

```
$ ./split.py conll2000.tag train.tag dev.tag test.tag
usage: split.py [-h] --seed SEED input train dev test
split.py: error: the following arguments are required: --seed
```

3. Run your script over the input file producing `train.tag`, `dev.tag`, and `test.tag` using a fixed seed. Then, compute their SHA-256 checksum using a tool like the `shasum` command-line tool:

```
$ shasum -a256 train.tag dev.tag test.tag
```

Then, run the script over the input file *again*, using the same seed as the previous time, and confirm that the SHA-256 checksums have not changed.

## Hints

- Read the [random docs](#) before you get started.
- There are in fact two ways to accomplish this, either fine:
  - seeding the global random number generator.
  - create a `random.Random` object, seed it, and use it for shuffling.

## Stretch goals

(Recall that these are optional, not required.)

1. Print out the number of sentences and tokens in each of the output files in some visually-appealing format using f-strings.
2. Do #1, but instead of printing it out, use [Python's built-in logger](#) and log the information at `INFO` level. This information will not show up when you run the command unless you set the log level to `INFO` or lower; do so inside the `if __name__ == "__main__":` block, immediately before you parse the command-line arguments, or allow the user to enable a verbose mode via an optional, boolean `-v` flag.
3. Install the `black` code reformatter and reflow your code using `black -l79`.
4. Install the `flake8` linter and make your code pass its checks.
5. Add type signatures to your functions. Then install the `mypy` static type checking tool and make your code pass its checks.

## What to turn in

1. Your final `./split.py` as per Step 3.
2. A one-page write-up, in PDF form, detailing any challenges you experienced and how you dealt with them.

## References

Tjong Kim San, E. F. and Buchholz, S. 2000. Introduction to the CoNLL-2000 shared task: chunking. In *Fourth Conference on Computational Natural Language Learning and the Second Learning Language in Logic Workshop*, pages 127-132.